




Written in English for clarity — some technical terms don't translate naturally yet

Simple, but sure

Fast, Free, Local, No Code: Your Specialized Legal LLM is Ready

Bypass cloud fees and Python code. Load the model file directly into any desktop app (LM Studio, Ollama) and start chatting instantly




Afza Azzindani



The Mission of Accessible Legal AI

2/6

- **The Challenge:** Specialized LLMs (like our fine-tuned legal model) often remain trapped behind high cloud costs, complex Python environments, and the need for dedicated, powerful GPU hardware. This limits adoption by individual practitioners and solo researchers.
- **The Goal:** To ensure our structured legal reasoning model is globally accessible—removing technical and financial barriers while guaranteeing user privacy.
- **The Solution:** We quantized the model into the GPT-Generated Unified Format (GGUF). This makes the specialized knowledge base portable, runnable on standard CPU/GPU desktop machines, and entirely code-free to deploy.
- **Action:** Swipe to see the tools that make private, local AI a reality today.




Afza Azzindani

The Deployment Trade-Off (Cloud vs. Local)



Feature	Cloud Deployment	Quantized Local Deployment (GGUF)
Compute Scale	High/Infinite: Access to enterprise-grade GPUs/TPUs for massive tasks	Low/Consumer: Optimized to run on standard CPU or consumer GPU (e.g., M-series or RTX cards)
Cost Model	High Runtime Cost: Pay-per-token or subscription model; costs scale with usage	Zero Runtime Cost: Initial file download is the only "cost" (leveraging free tools)
Data Privacy	Low Privacy: Queries must be sent across the internet to the provider's server	High Privacy: Queries and responses remain entirely on the user's local machine
Accessibility	Requires complex API integration , keys, or high-speed internet dependency	High Accessibility: Plug-and-play setup via desktop apps (No Python code required)
Use Case	Batch processing, large-scale inference, and training.	Interactive Chat, Few-shot Testing, and private consultation where security is paramount.

- For **mass accessibility** and **individual user security** in sensitive domains like law, the GGUF model offers a vital alternative to the powerful but costly cloud model.




Quantization for Accessibility



The GGUF model is created through quantization, a process that allows for **efficient local deployment** but involves a strategic trade-off

Feature	Full Precision Model (Original FP16/32)	Quantized Model (GGUF/INT4-8)
Data Representation	Uses 16 or 32 bits per weight (highest detail/precision).	Uses 4 to 8 bits per weight (compact detail).
File Size	Massive (e.g., ~16 GB for a 8B model)	Small (e.g., ~6 GB for the same model)
Hardware Requirement	Requires dedicated VRAM/GPU (limits local access)	Runs efficiently on CPU or lower VRAM/GPU (maximum accessibility)
Inference Speed	Fastest on high-end GPUs	Fast on Consumer Hardware : Optimized for desktop parallelism.
Accuracy/Quality	100% of fine-tuned precision. Used for final R&D and benchmarking.	~97 to 99% of original precision

- **The Justification:** The GGUF format strategically accepts a **minimal accuracy drop** to achieve a **massive gain in deployment velocity**, cost reduction, and accessibility for the end-user.



Plug-and-Play Deployment



The GGUF format **enables code-free**, cross-platform access. Follow these steps to run the model:

1. **Download:** Get the GGUF file from the model's repository page
2. **Load:** Open your preferred desktop runner and load the GGUF file directly (e.g., LM Studio, Ollama, Jan, or KoboldCPP)
3. **Customize System Prompt** (CRITICAL STEP): Paste the following system instruction into the configuration window:

Anda adalah asisten AI yang ahli di bidang hukum Indonesia. Anda dapat membantu konsultasi hukum, menjawab pertanyaan, dan memberikan analisis berdasarkan peraturan perundang-undangan yang relevan.

Untuk setiap respons, Anda harus berfikir dan menjawab dengan Bahasa Indonesia, serta gunakan format:

<think>

...

</think>

Tuliskan jawaban akhir secara jelas, ringkas, profesional, dan berempati jika diperlukan. Gunakan bahasa hukum yang mudah dipahami. Sertakan referensi hukum Indonesia yang relevan. Selalu rekomendasikan konsultasi dengan ahli hukum untuk keputusan final.

- **The Result:** The model will now run locally, privately, and generate the structured, cited legal reasoning developed during the GRPO fine-tuning
- **Repository:** [Azzindani/Deepseek ID Legal Preview GGUF](#)
- **Acknowledgements:**
 1. [Qwen team](#)
 2. [Deepseek team](#)
 3. [GRPO research paper](#)



This project validates that specialized AI can be aggressively quantized (GGUF) without sacrificing core utility.

Interested in leveraging local compute for high-stakes, specialized projects?

Let's talk—your insights and experience are essential for this research



Afza Azzindani

