



Written in English for clarity — some technical terms don't translate naturally yet

**Start today.
tomorrow, it's done.**

From PDFs to Power: 3.5 Million Legal Articles Parsed

Introducing the Indonesian
Regulation Parsed Dataset ~250,000
regulations now structured at the
atomic article level.




Afza Azzindani



The Data Structuring Imperative

2/6


- **Motivation:** My drive is to make legal information trustworthy and accessible. The key is structuring raw, messy data into a usable format for any future tool.
- **The Problem:** The initial ~250K PDF collection lacked the granularity for precise citation; standard chunking breaks the legal context.
- **The Solution:** I built a pipeline to parse text strictly based on the original document's Article, Clause, and Section structure.
- **Disclaimer:** This dataset is built exclusively for R&D using publicly accessible documents. It is provided without warranty regarding the completeness, accuracy, or legal validity of the content.



Trade-Off: Speed & Cost vs. Cleanliness

Metric	Manual Structuring	Automated Parsing
Workflow	Open PDF → Read → Copy/Paste each Article → Label (Slow, subjective, error-prone)	Extracts raw text → Parses article boundaries → Labels based on predefined rules (Fast, consistent, scalable)
Time Consumption	~60,000 total hours. Prohibitively slow even when scaled	~30 hours of distributed processing time
Cost	Massive Financial Burden over months/years	Near-Zero Cost by leveraging free cloud runtime/storage
Downside	Fatigue, mislabeling are possible but generally low	Typos, minor formatting errors, or missed article breaks due to inconsistent original PDF layouts
Margin of Error	High accuracy (low error)	The goal is to keep errors low enough that they can be filtered out in post-processing without corrupting the 3.5 Million rows


- The Point: In R&D, I must choose **Strategic Velocity over Surgical Precision**. The benefit of immediately generating 3.5 Million structured articles outweighs the cost of manually eliminating every minor error, as "**what's usable now is far more valuable than what's flawless but locked away for years**" This project trades minimal automated error for maximum accessibility and scale



Strategic Efficiency & Portability

Constraint	Local Machine (PC, Laptop)	Cloud Strategy (Google Colab & Hugging Face)
Setup & Environment	Manual installation of parsers/dependencies; prone to version conflicts	Portable Environment: is consistent, containerized, and immediately available via browser
Physical Location	Task is bound to the physical location of the PC (no movement allowed during processing)	Full Portability: I could run & monitor the process from any device, anywhere via the web browser
Multitasking	PC is sluggish or unusable for other tasks during intensive processing	PC Free: Local machine's CPU/RAM are available, allowing me to work on other projects while processing runs in the cloud

- **Final Architecture:** The workload was split across **6 Google Colab instance**, maximizing speed and fault tolerance, with output **pushed directly to Hugging Face**.



High-Value Use Cases & Next Steps



This highly granular dataset is engineered for precision-focused legal AI and open R&D:

- **Over 3.5 Million rows** (articles/clauses) parsed from **250K documents**.
- **Time Cost: ~30 hours** of distributed data extraction & parsing
- Content includes the **full hierarchy of Indonesian law**, such as UU, PP, Perpres, Permen, Pergub, and Perda (main regulations and explanations).
- **Repository:** [Azzindani/ID REG Parsed](#)
- **Key Applications:**
 1. Training LLMs on legal text with **natural article boundaries**.
 2. Building **RAG systems** that retrieve the exact legal clause for compliance tools.
 3. Enabling **article-level classification** and fine-grained analysis.
- **Acknowledgements:**
 1. [Hugging Face](#)
 2. [Google Colab](#)





This project proves that effective R&D at massive scale is a matter of strategic resourcefulness, not budget size.

Is there a better way to increase the extraction and parsing accuracy while keeping the process cost-efficient and scalable?

Share your ideas for maximizing the utility of this open resource.



Afza Azzindani